

# Diabetes Prediction using Logistic Regression

( Machine Learning )



<https://www.linkedin.com/in/nibedita-sarkar/>



# logisticregression-1

June 8, 2023

## 1 Health Parameter Analysis and Diabetes Prediction using Logistic Regression

Logistic regression is a useful model for predicting binary outcomes, where there are only two possible classes. It is commonly used because it provides interpretable results, allows for estimating probabilities, and is computationally efficient. Logistic regression has fewer assumptions compared to other models, can handle non-linear relationships, and is robust to outliers. It also supports regularization techniques to prevent overfitting. Logistic regression is a well-studied and established model in statistics and machine learning. However, it may not be suitable for all classification problems, especially those with highly nonlinear relationships.

### 1.1 Define the Problem :

The problem is to create a logistic regression model based on the provided dataset that can predict the outcome of diabetes based on health parameters. The dataset contains various columns, including the outcome column, which represents whether a person has diabetes or not. The goal is to train the model using the training set and evaluate its performance on the testing set using the confusion matrix and accuracy score.

### 1.2 Importing necessary libraries :

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix
import warnings
warnings.filterwarnings('ignore')
```

### 1.3 Build the dataset:

To build the dataset, you need to perform the following steps:

**a) Load the dataset using pandas:** Use the pandas library to load the dataset from the 'diabetes.csv' file

```
[2]: data = pd.read_csv("diabetes.csv")
data.head(5)
```

```
[2]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

**b) Exploring the dataset** Exploring the dataset is crucial for understanding its structure, identifying missing values, data distribution, correlations, and outliers, as well as for making informed decisions regarding data preprocessing and feature selection. It provides insights that guide data analysis and model building.

```
[3]: data.shape # Checking number of rows and coulmn in the dataset
```

```
[3]: (768, 9)
```

```
[4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                              768 non-null    int64
2   BloodPressure                        768 non-null    int64
3   SkinThickness                       768 non-null    int64
4   Insulin                             768 non-null    int64
5   BMI                                 768 non-null    float64
6   DiabetesPedigreeFunction             768 non-null    float64
7   Age                                 768 non-null    int64
8   Outcome                             768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

`data.info()` - Checking the information about the dataset, including the number of entries, the number of columns, column names, data types of each column, and any missing values. It will help in understanding the structure and properties of the dataset.

```
[5]: data.describe()
```

```
[5]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

data.describe() - Checks the descriptive statistics for the numerical columns in the dataset. The statistics include count, mean, standard deviation, minimum value, 25th percentile (Q1), median (50th percentile or Q2), 75th percentile (Q3), and maximum value. It provides a summary of the central tendency, spread, and distribution of the numerical data.

```
[6]: data.isnull().sum() # Checking if the dataset contains any null values.
```

```
[6]: Pregnancies      0
      Glucose         0
      BloodPressure   0
      SkinThickness   0
      Insulin         0
      BMI             0
      DiabetesPedigreeFunction  0
      Age             0
      Outcome         0
      dtype: int64
```

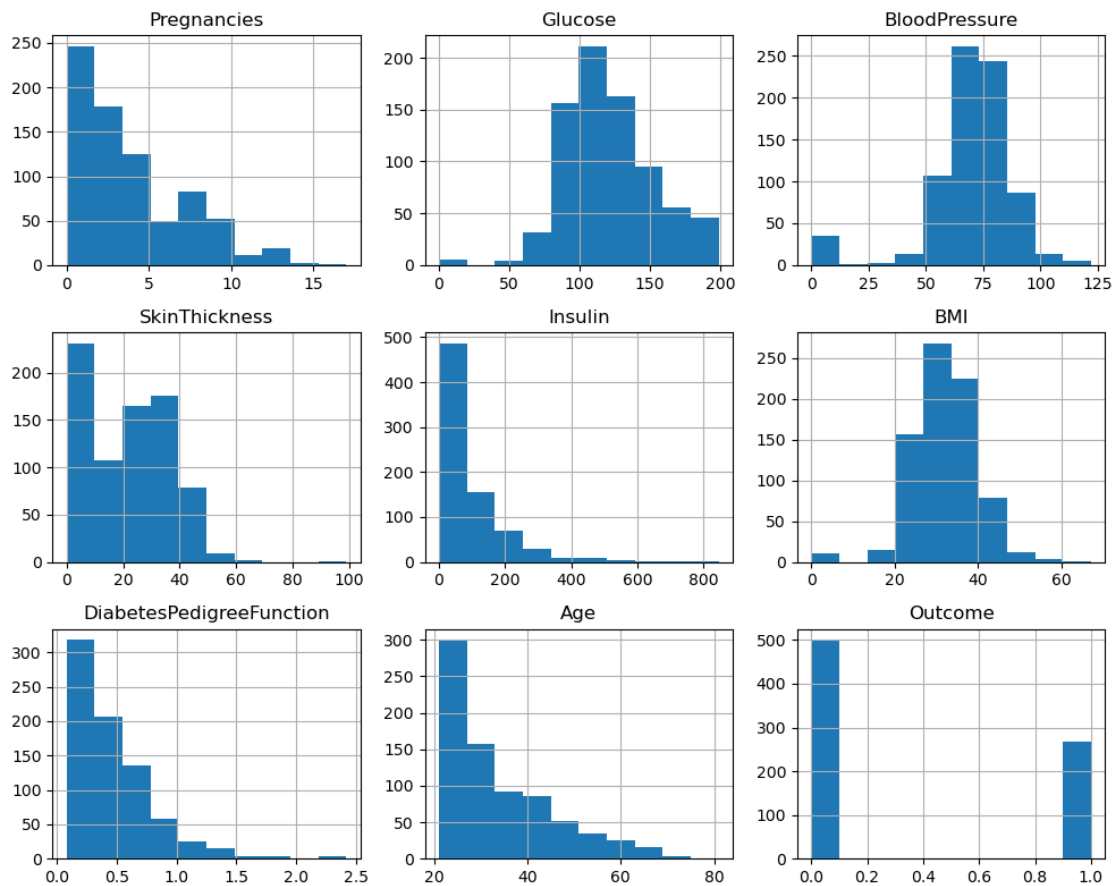
```
[7]: num_duplicates = data.duplicated().sum() # Checking if the dataset contains any
      ↪ duplicate values

      if num_duplicates > 0:
          print(f"The dataset contains {num_duplicates} duplicate values")
          data = data.drop_duplicates
          print("Number of duplicate values after dropping:", num_duplicates)
```

```
else:
    print("The dataset doesn't contain any duplicate values.")
```

The dataset doesn't contain any duplicate values.

```
[8]: data.hist(figsize=(10, 8)) # Checking Data Distribution
plt.tight_layout()
plt.show()
```



c) **Extract data from the outcome column as a variable named Y:** Extract the values from the 'outcome' column and assign them to a variable called Y.

```
[9]: X = data.iloc[:, :-1]
X.head(5)
```

```
[9]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  \
0             6      148             72             35         0  33.6
1             1       85             66             29         0  26.6
2             8      183             64              0         0  23.3
```

3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age
0	0.627	50
1	0.351	31
2	0.672	32
3	0.167	21
4	2.288	33

d) **Extract data from every column except the outcome column as a variable named X:** Extract the data from all columns except the 'outcome' column and assign them to a variable called X.

```
[10]: Y = data.iloc[:, -1]
      Y.head(5)
```

```
[10]: 0    1
      1    0
      2    1
      3    0
      4    1
      Name: Outcome, dtype: int64
```

e) **Divide the dataset into two parts for training and testing:** Split the dataset into a training set and a testing set in a 70% - 30% proportion. This will be used to train the model on the training set and evaluate its performance on the testing set.

```
[11]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.30,
      ↪ random_state = 51)
```

```
[12]: print(X_train.shape)
      print(X_test.shape)
      print(Y_train.shape)
      print(Y_test.shape)
```

```
(537, 8)
(231, 8)
(537,)
(231,)
```

## 1.4 Train the model:

```
[13]: logistic = LogisticRegression()
      logistic.fit(X_train, Y_train)
```

```
[13]: LogisticRegression()
```

## 1.5 Evaluate the model :

```
[14]: Y_predict = logistic.predict(X_test)
      print("Y_predict:\n",Y_predict)
```

```
Y_predict:
[0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 1 1 0 0 1 0 0 0 0 1 1 0 1 0 1 0 0 1 0 0 0
 0 0 0 0 0 0 1 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 1 0 0 0 1 0 0 1
 0 1 0 1 1 0 0 0 1 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 1 1 0 1 0 0
 1 0 0 0 1 1 1 0 0 0 0 0 0 1 0 1 0 0 1 0 0 1 0 0 1 0 1 0 1 0 0 0 0 0 0 0 0 0
 0 0 0 1 1 0 0 0 0 0 0 0 0 1 0 1 1 0 1 0 1 1 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 1
 1 1 0 0 1 0 0 0 1 0 0 0 0 0 0 1 1 0 0 1 0 0 0 0 1 0 1 1 0 0 1 0 0 1 0 0 1 1 1 0
 0 0 0 0 1 0 1 0 0]
```

```
[15]: print("Y_test:\n",Y_test)
```

```
Y_test:
737    0
505    0
296    1
711    0
329    0
..
405    0
315    0
131    1
364    0
322    1
Name: Outcome, Length: 231, dtype: int64
```

```
[16]: score = accuracy_score(Y_test, Y_predict)
      print("Accuracy Score: ",score * 100)
```

```
Accuracy Score: 79.22077922077922
```

```
[17]: confusion_matrix = confusion_matrix(Y_test, Y_predict)
      print("Confusion Matrix : \n",confusion_matrix)
```

```
Confusion Matrix :
[[131  11]
 [ 37  52]]
```

## 1.6 Visually Understanding the performance of the model

```
[18]: # Create a figure and axis
      fig, ax = plt.subplots()

      # Plot the actual outcomes
```



```

ax.scatter(range(len(Y_test)), Y_test, color='blue', label='Actual Outcome')

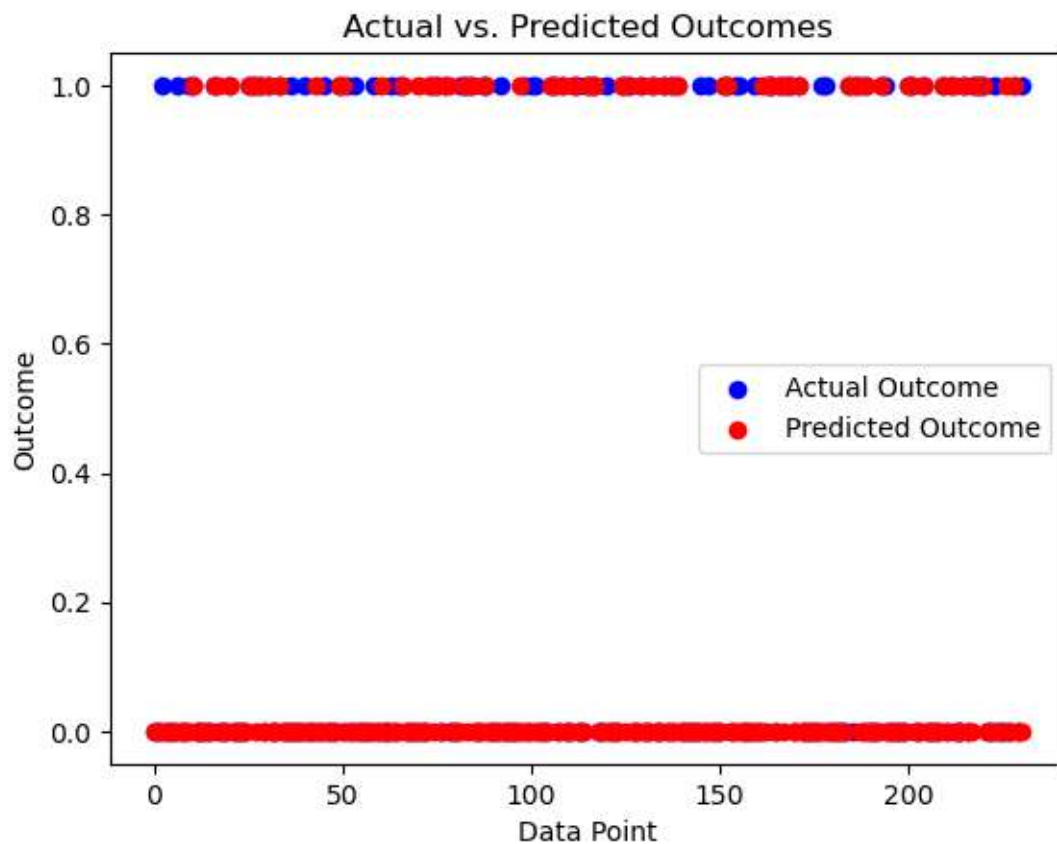
# Plot the predicted outcomes
ax.scatter(range(len(Y_predict)), Y_predict, color='red', label='Predicted Outcome')

# Set axis labels and title
ax.set_xlabel('Data Point')
ax.set_ylabel('Outcome')
ax.set_title('Actual vs. Predicted Outcomes')

# Add a legend
ax.legend()

# Show the plot
plt.show()

```





## **1.7 Use the model:**

Once the model is trained and evaluated, you can use it to make predictions on new, unseen data. This can be done by providing new input values to the model and using the predict function to obtain the predicted outcome.